

Introduction to Machine Learning & Deep Learning Concepts

July 2021

Zafri Baharuddin

Based on work by
Dickson Neoh
Justin Johnson



intel®

ITXOTIC

Deep Learning for Computer Vision

Building artificial systems that
process,
perceive, and
reason about visual data

Computer Vision is everywhere



Left to right:
[Image by Intel](#) is
copyrighted
[Image](#) from [Unsplash](#)
[Image](#) from [Unsplash](#)
[Image](#) by [Nasa](#) is public



Left to right:
[Image](#) by [Nasa](#) is public
[Image](#) from [Unsplash](#)
[Image](#) from [Unsplash](#)
[Image](#) from [Unsplash](#)



Left to right:
[Image](#) from [Unsplash](#)
[Image](#) from [Unsplash](#)
[Image](#) from [Unsplash](#)
Image copyrighted by
[ExtremeTech](#)

Deep Learning for Computer Vision

Building artificial systems that
learn from data and experience

Deep Learning for Computer Vision

Hierarchical learning algorithms
with many “layers”,
(very) loosely inspired by the brain

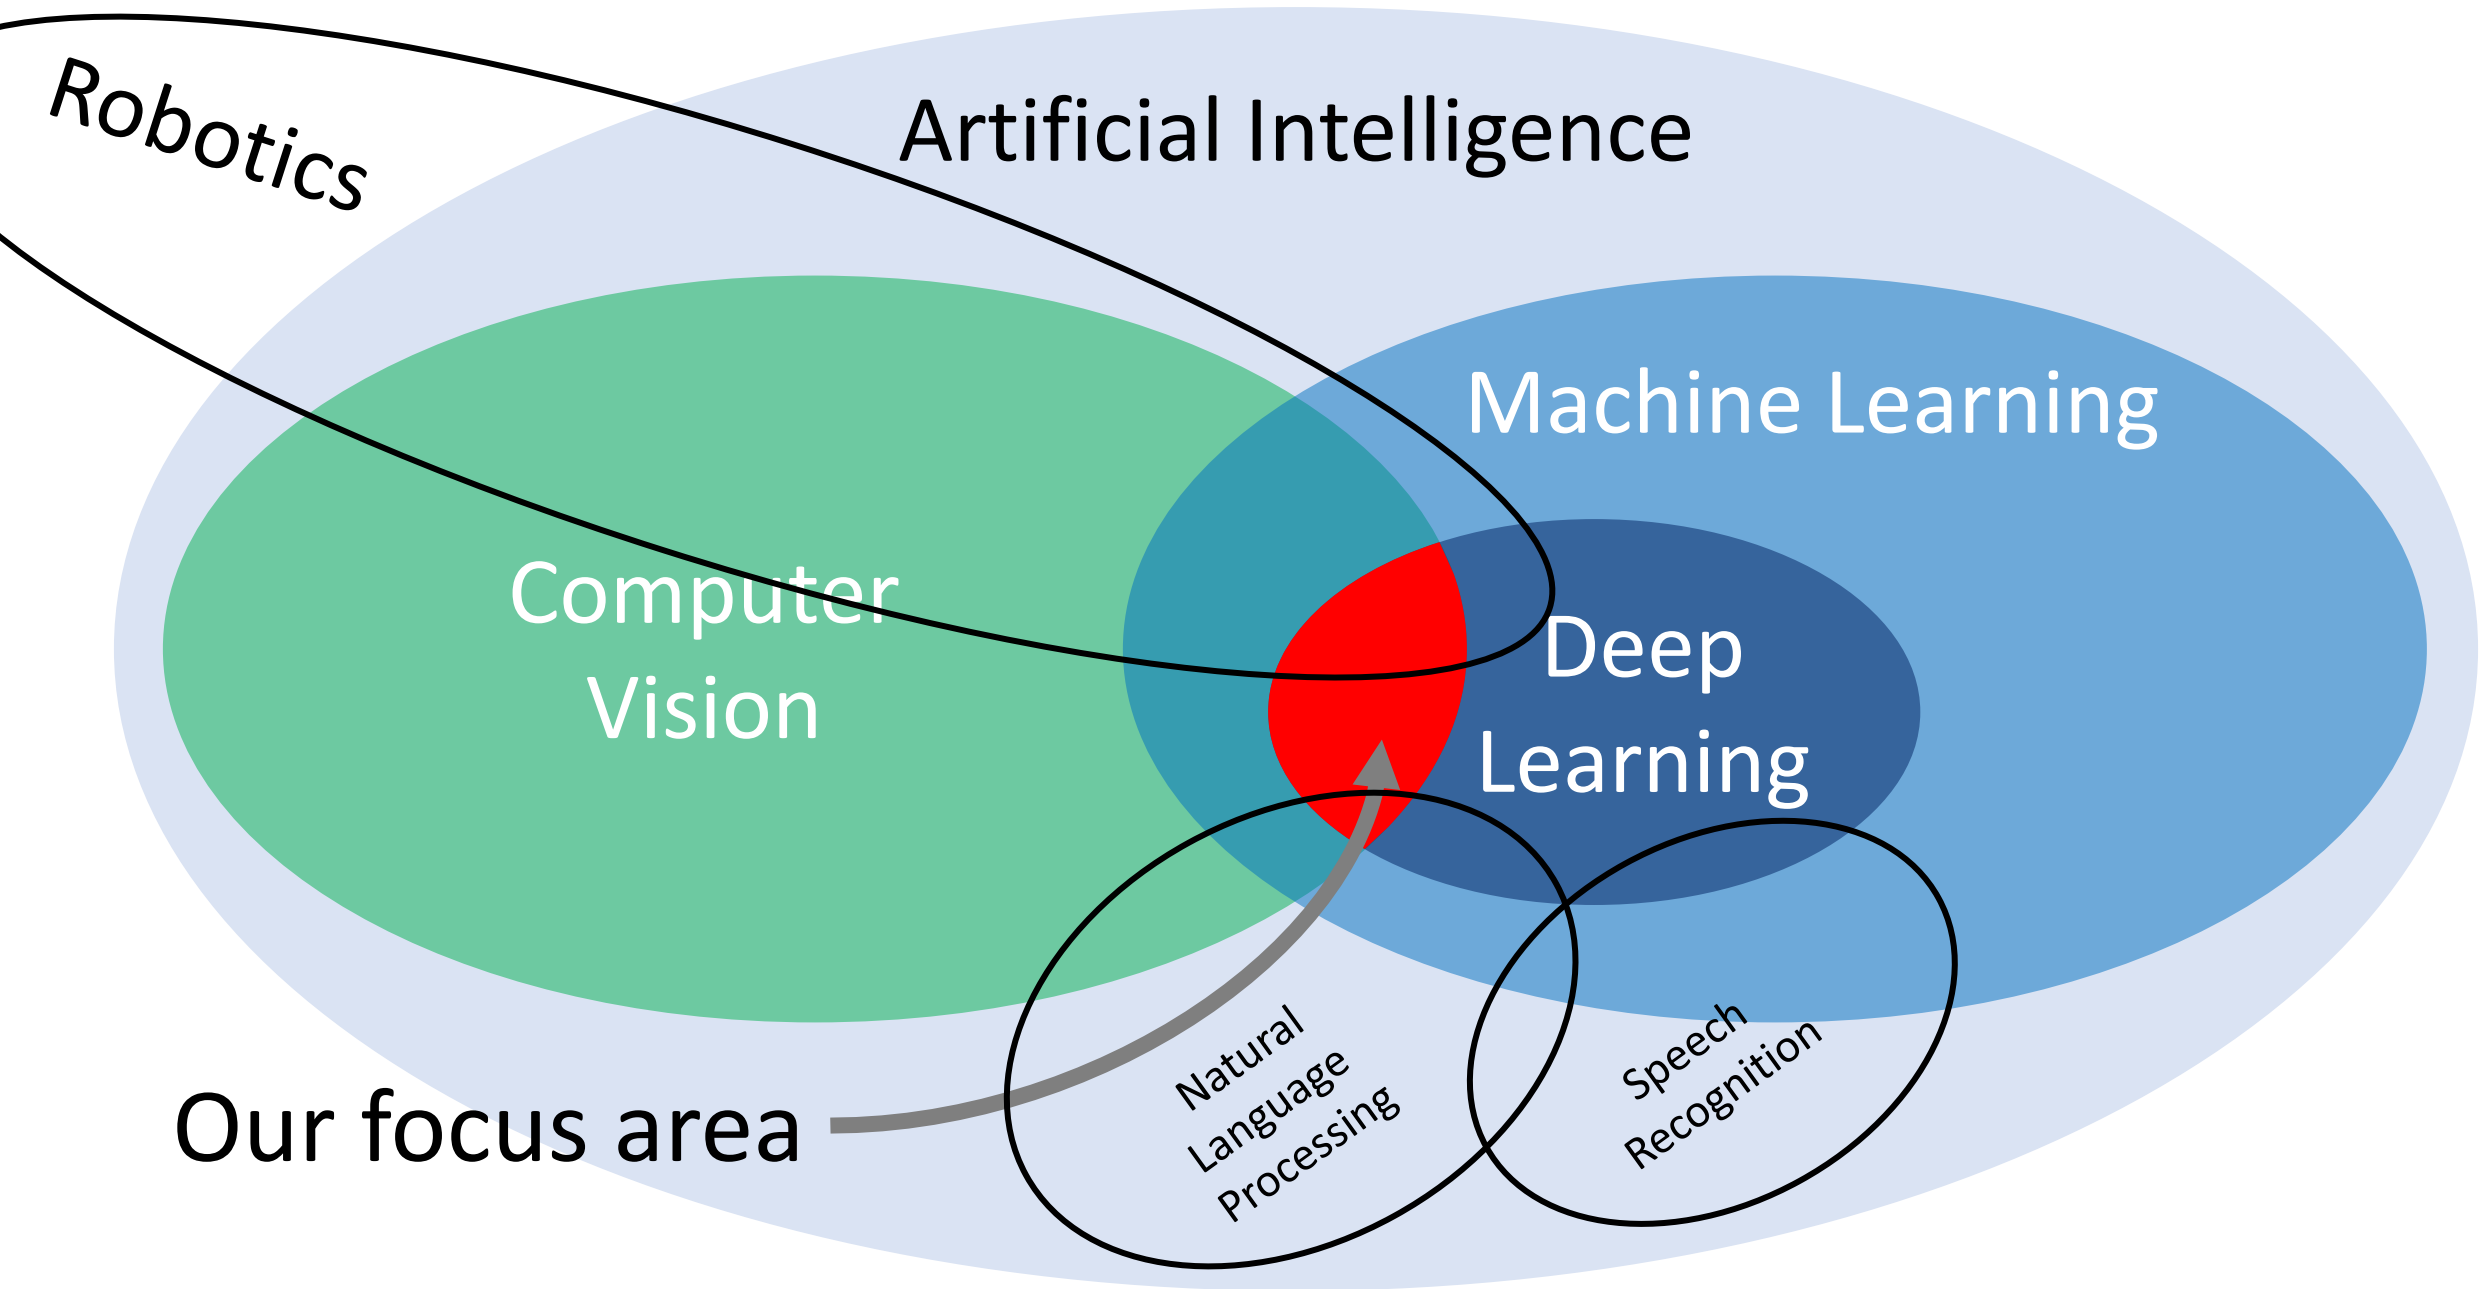


Image Classification: Challenges

- Change view angle



Image Classification: Challenges

- Variation within class




Image Classification: Challenges

- Sub-classes (different cat breeds)



Maine Coon



British Shorthair



Grumpy Cat
(Mixed breed)
2012-2019

Image Classification: Challenges

- Background blending / camouflage




Image Classification: Challenges

- Lighting change



Image Classification: Challenges

- Deformation



Image Classification: Challenges

- Occlusion



Image Classification: Building Blocks for other tasks.

- Object detection



Output: bounding boxes and class

Image Classification: Building Blocks for other tasks.

- Object segmentation



Output: segments and class

Image Classification: Building Blocks for other tasks.

- Image captioning



Output: description: There are two cars and a bike on the road

Motivation: Classification – Basic Human task

Input: image



Resolution:
800 x 600 x 3 (RGB)



Output:

cat
bird
deer
dog
truck

Motivation: Classification – Basic AI task

Input: image




Resolution:
800 x 600 x 3 (RGB)

AI
Magic
Box

Output: Probability:

cat	0.82
bird	0.02
deer	0.04
dog	0.10
truck	0.02



Numbers

240	200	150	100	50	20	15
200	200	150	100	50	20	15
150	150	150	100	50	20	20
100	100	100	100	50	50	50
50	50	50	50	100	100	100
20	20	20	50	100	150	150
15	15	20	50	100	150	200



Black Box






Numbers

0
1

Numbers

Numbers



Numbers

Numbers

Deep
learning
for
humans



Black Box



인간을위
한 딥 러닝

Numbers




Black Box




Numbers

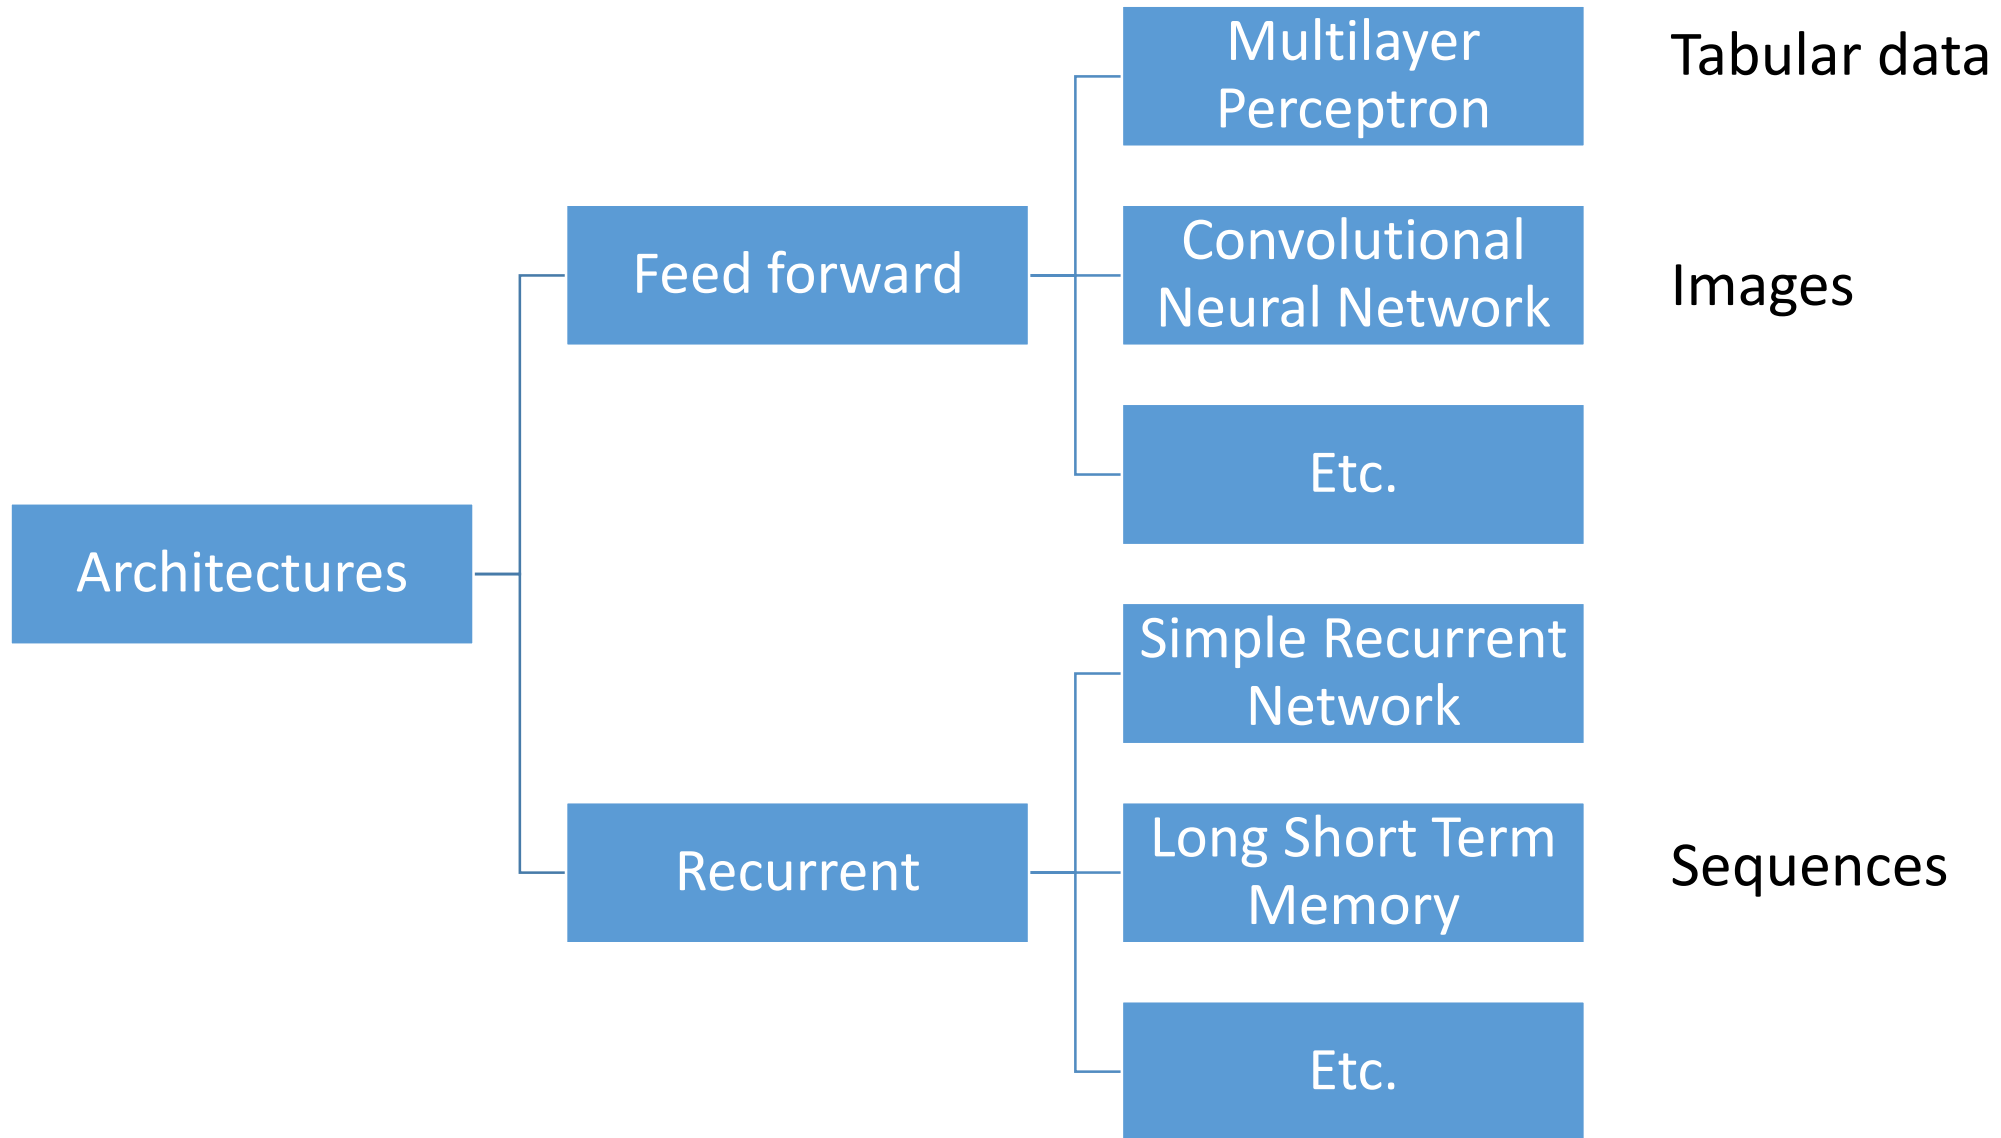
인간을위
한 딥 러닝




Weight

Neuron







Feedforward neural network



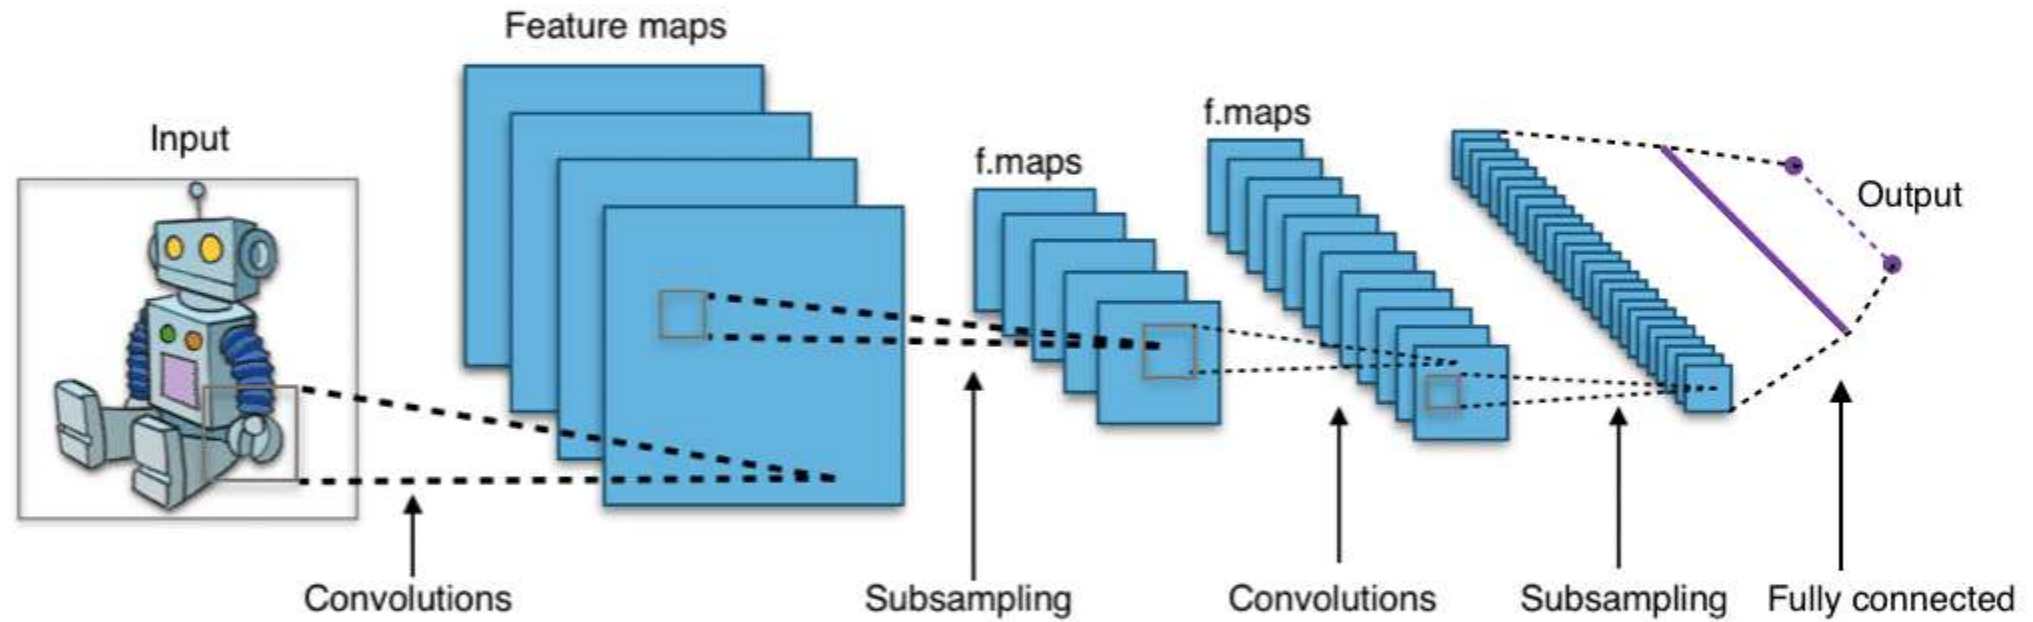
Recurrent neural network



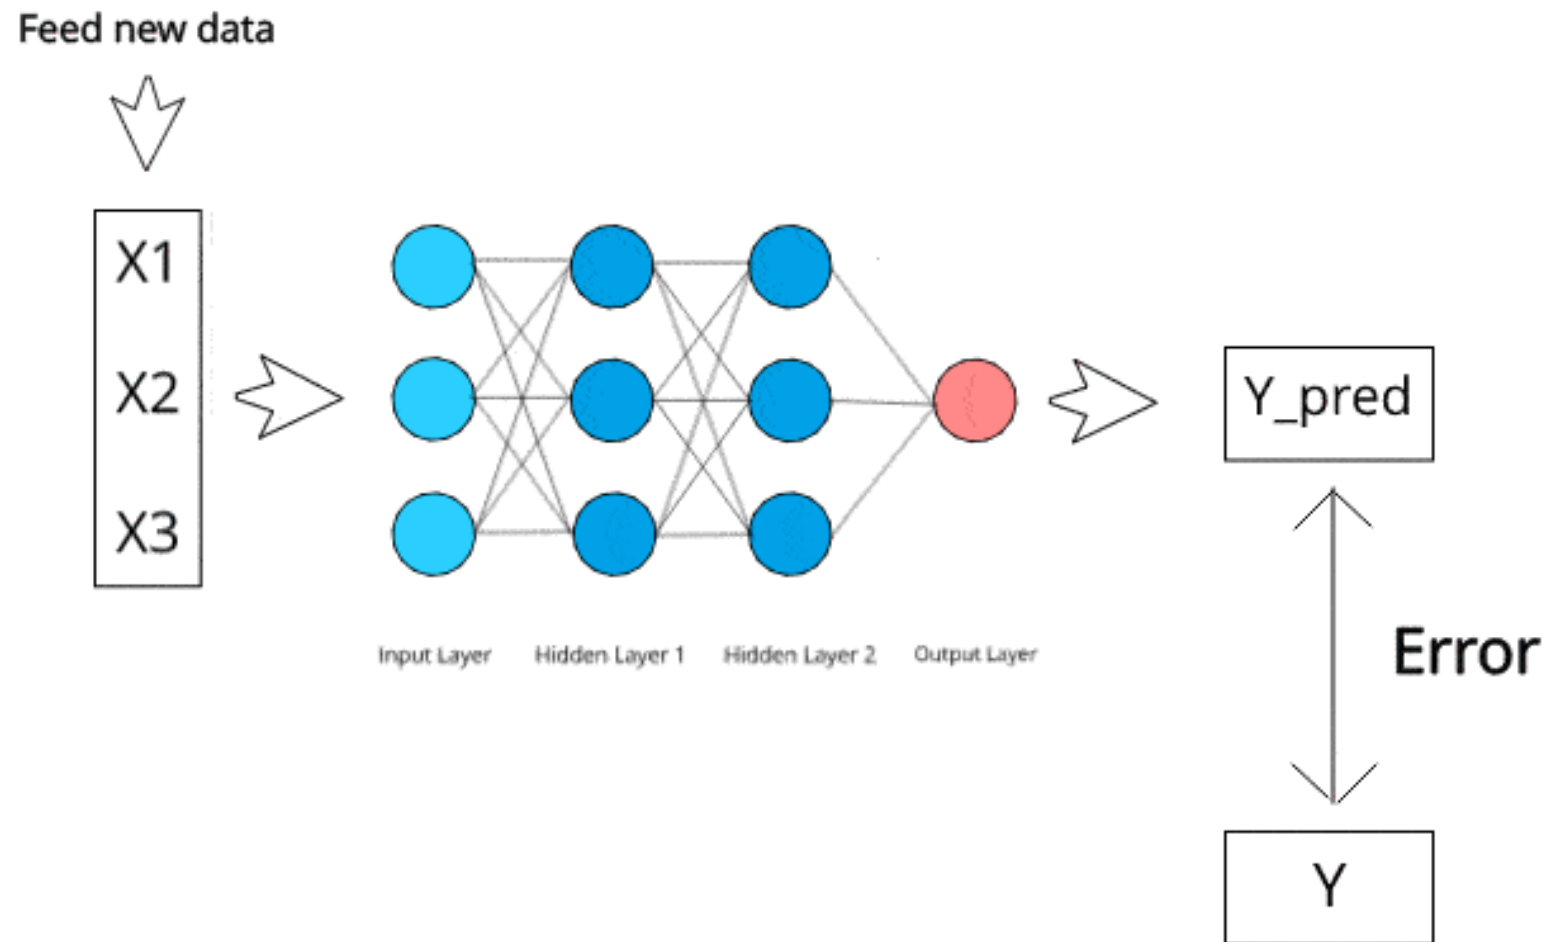
MLP



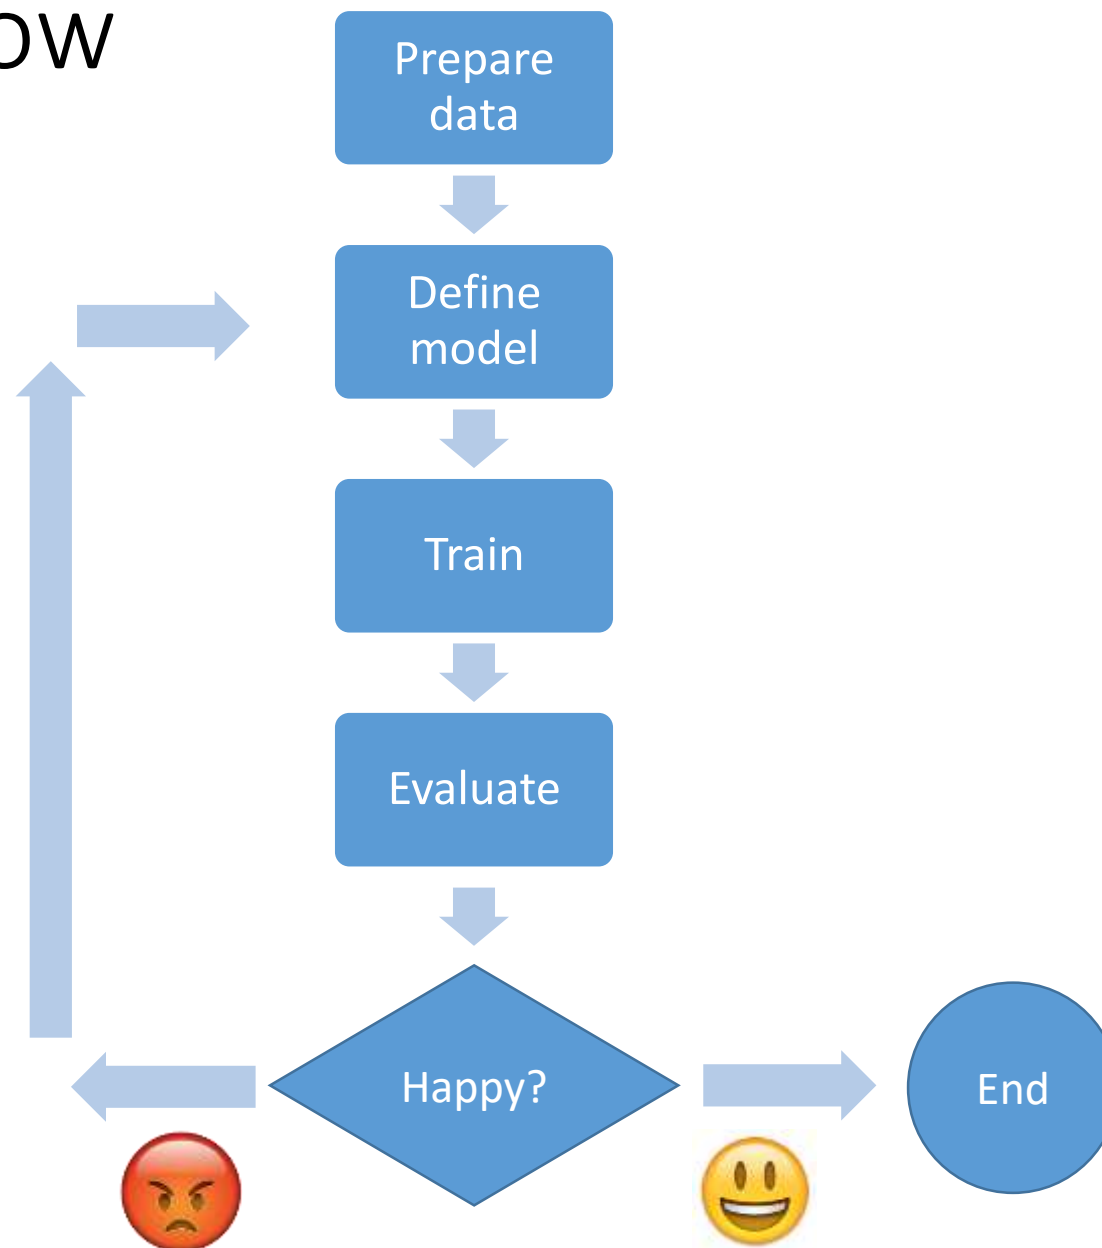
CNN



Learning with Backpropagation



Typical DL workflow




How

- Learn by doing.
- Minimal theory/math to get started with Tensorflow

What

- Jupyter Notebook
- Google Colab (or go to colab.research.google.com)
- Interactive demo

Lab 1: MNIST with MLP




MNIST


- Modified National Institute of Standards and Technology
- The Hello World of Deep Learning
- 10 classes (0-9)
- Image size: 28 x 28 pixels
- SOTA: 99.84%
- Human: 99.80%



One hot encoding



MLP




MLP: 97.0 ++ %

- 97+% accuracy with < 20 lines of codes
- Open [MNIST with MLP](#)

```
1. import numpy as np
2. from tensorflow import keras
3. num_classes = 10
4. input_shape = (28, 28, 1)
5. (x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()
6. x_train = x_train.astype("float32") / 255
7. x_test = x_test.astype("float32") / 255
8. x_train = np.expand_dims(x_train, -1)
9. x_test = np.expand_dims(x_test, -1)
10. y_train = keras.utils.to_categorical(y_train, num_classes)
11. y_test = keras.utils.to_categorical(y_test, num_classes)
12. model = keras.Sequential([
13.     keras.Input(shape=input_shape),
14.     keras.layers.Flatten(),
15.     keras.layers.Dense(128, activation='relu'),
16.     keras.layers.Dense(num_classes, activation="softmax"),
17. ])
18. model.compile(loss="categorical_crossentropy", optimizer="adam", metrics=["accuracy"])
19. model.fit(x_train, y_train, batch_size=128, epochs=5, validation_split=0.1)
```



Hyperparameters

```
model = keras.Sequential(  
    [  
        keras.Input(shape=input_shape),  
        layers.Flatten(),  
        layers.Dense(128, activation='relu'),  
        layers.Dense(num_classes, activation="softmax"),  
    ]  
)  
  
model.summary()
```




784 neurons

Lab 2: MNIST with CNN




Convolution Operation

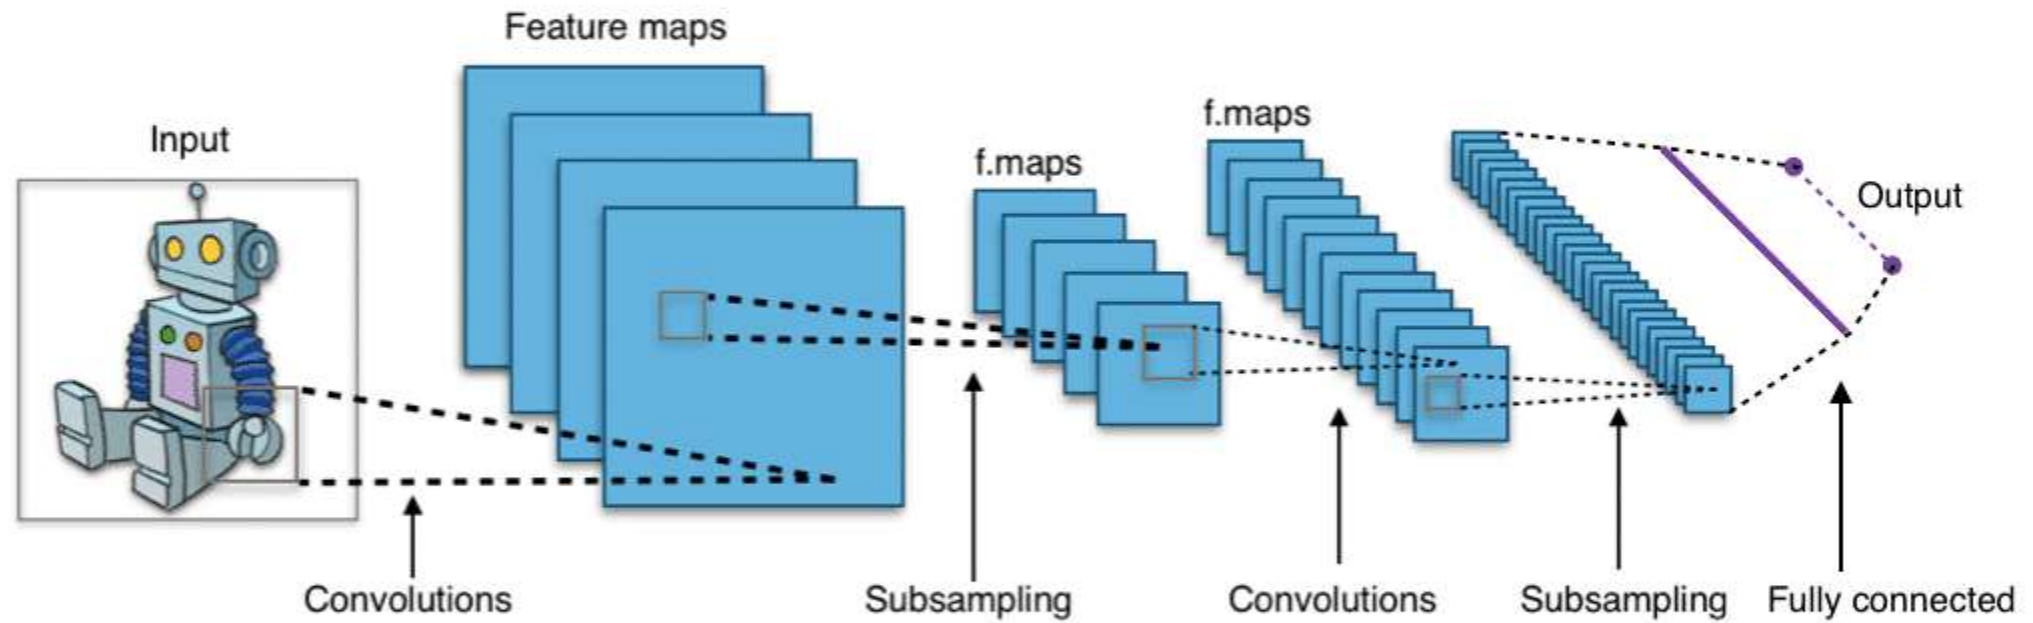


CNN: 98.0 ++ %

```
model = keras.Sequential(  
    [  
        keras.Input(shape=input_shape),  
        layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),  
        layers.Flatten(),  
        layers.Dense(num_classes, activation="softmax"),  
    ]  
)
```




CNN




CNN Lenet: 99.0 ++ %

Open [MNIST with Lenet](#)

```
model = keras.Sequential(  
    [  
        keras.Input(shape=input_shape),  
        layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),  
        layers.MaxPooling2D(pool_size=(2, 2)),  
        layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),  
        layers.MaxPooling2D(pool_size=(2, 2)),  
        layers.Flatten(),  
        layers.Dense(64, activation='relu'),  
        layers.Dense(num_classes, activation="softmax"),  
    ]  
)
```



Max Pooling Operation



Feature maps learned from data




Image Features: Color Histogram



Ignores texture,
spatial positions





Image Features: Histogram of Oriented Gradients (HoG)




Weak edges
Strong diagonal
edges



Edges in all
directions

Captures
texture and
position,
robust to
small image
changes



1. Compute edge direction / strength at each pixel
2. Divide image into 8x8 regions
3. Within each region compute a histogram of edge directions weighted by edge strength

Example: 320x240 image gets divided into 40x30 bins; 8 directions per bin; feature vector has $30 \times 40 \times 9 = 10,800$ numbers

Lowe, "Object recognition from local scale-invariant features", ICCV 1999
Dalal and Triggs, "Histograms of oriented gradients for human detection," CVPR 2005

Image Features: Bag of Words (Data-Driven)


Step 1: Build codebook




Extract random
patches
from all images



Cluster patches to
form “codebook”
of “visual words”



Step 2: Encode images



Fei-Fei and Perona, “A bayesian hierarchical model for learning natural scene categories”, CVPR 2005

Image Features

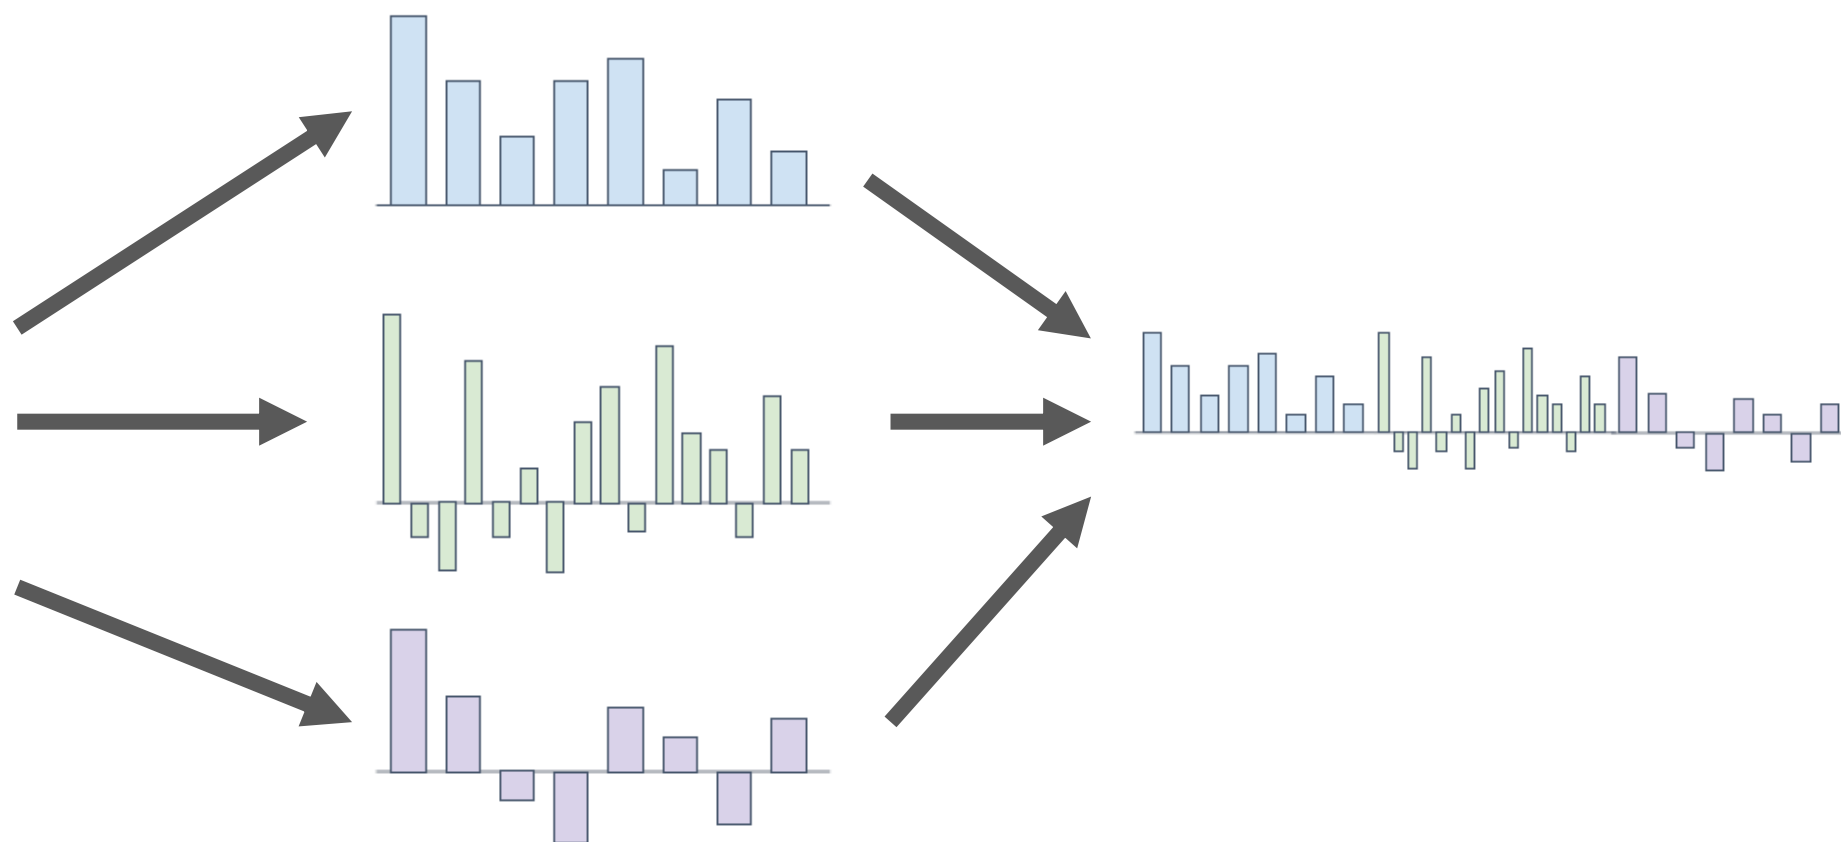


Image Features



Feature Extraction



f




10 numbers giving
scores for classes




training

Image Features vs Neural Networks



10 numbers giving scores for classes



Krizhevsky, Sutskever, and Hinton, "Imagenet classification with deep convolutional neural networks", NIPS 2012.
Figure copyright Krizhevsky, Sutskever, and Hinton, 2012.

10 numbers giving scores for classes

Faces










Cars



Elephants



Chairs



Activation Function


ReLU is a good default choice
for most problems

- A function that defines the output of a neuron given the input

The function


$$\text{ReLU}(z) = \max(0, z)$$

is called “Rectified Linear Unit”




Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$


ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$




Non-linear curve with activation function


Linear




Nonlinear




Points not linearly
separable in original space




Points not linearly separable in original space




Feature transform:
 $h = Wx$








Feature transform:
 $h = Wx$




Consider a neural net hidden layer:
 $h = \text{ReLU}(Wx) = \max(0, Wx)$
Where x, h are both 2-dimensional



Feature transform:
 $h = \text{ReLU}(Wx)$





Consider a neural net hidden layer:
 $h = \text{ReLU}(Wx) = \max(0, Wx)$
Where x, h are both 2-dimensional



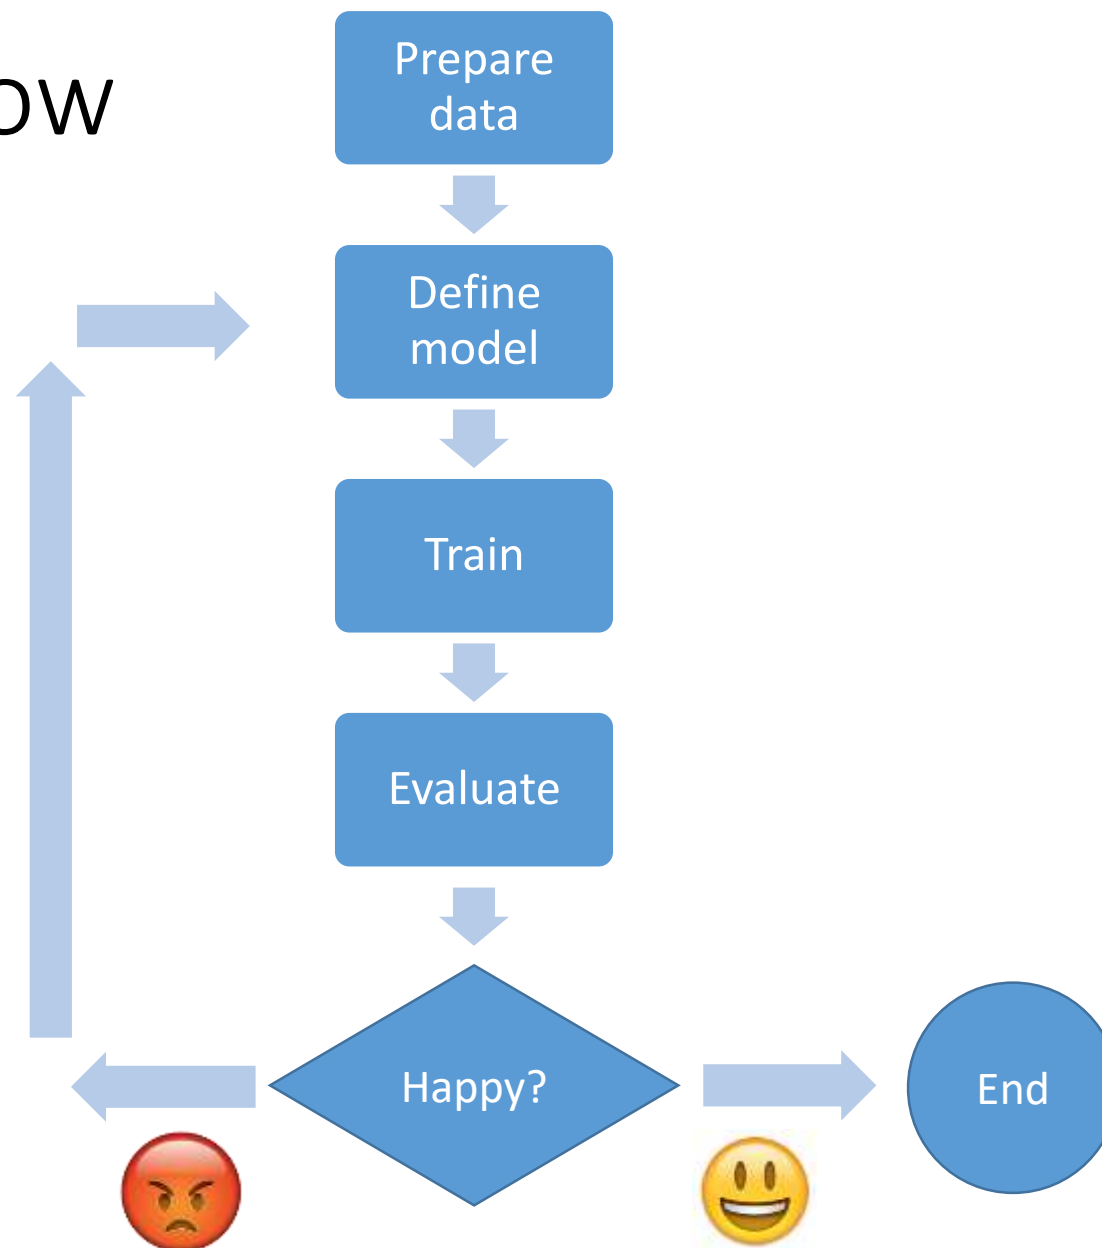
Linear classifier in feature space gives nonlinear classifier in original space

Feature transform:
 $h = \text{ReLU}(Wx)$



Points are linearly separable in features space!

Typical DL workflow




CHALLENGE

- Train the best classifier on the MNIST dataset
- Tune the following hyperparameters
 - # filters
 - Kernel size
 - # neurons
 - Activation function: 'relu', 'sigmoid', 'tanh'
 - # epochs
 - # layers


Overfitting

- Data Augmentation
- Use Dropout layer
- Transfer Learning


Learning curve




Overfitting



Underfit




Optimal




Overfit

Signs of overfitting

- High training accuracy, low validation/test accuracy




Normal



Overfit

Dropout Regularization

```
model = keras.Sequential(  
    [  
        keras.Input(shape=input_shape),  
        layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),  
        layers.MaxPooling2D(pool_size=(2, 2)),  
        layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),  
        layers.MaxPooling2D(pool_size=(2, 2)),  
        layers.Flatten(),  
        layers.Dense(64, activation='relu'),  
        layers.Dropout(0.2),  
        layers.Dense(num_classes, activation="softmax"),  
    ]  
)
```




Dropout randomly sets input units to 0, which helps prevent overfitting.

Data Augmentation

More data variation helps reduce overfitting

- Rotation
- Translation
- Flipping
- Zoom/Pan






Enlarge your Dataset

Data Augmentation

More data variation helps reduce overfitting

- Rotation
- Translation
- Flipping
- Zoom/Pan



Data Augmentation



Mirror + Rotation



Rotation

Data Augmentation



Mirror + Noise



Translation

CHALLENGE


- Implement regularization techniques.
- Tune hyperparameters to yield the best accuracy.
- Try on other datasets.

CIFAR10

Datasets

1. MNIST
2. CIFAR10
3. CIFAR100
4. ImageNet
5. MIT Places


airplane
automobile
bird
cat
deer
dog
frog
horse
ship
truck



50,000 training images
Each image is **32x32x3**

10,000 test images.

Datasets: Number of Training Pixels

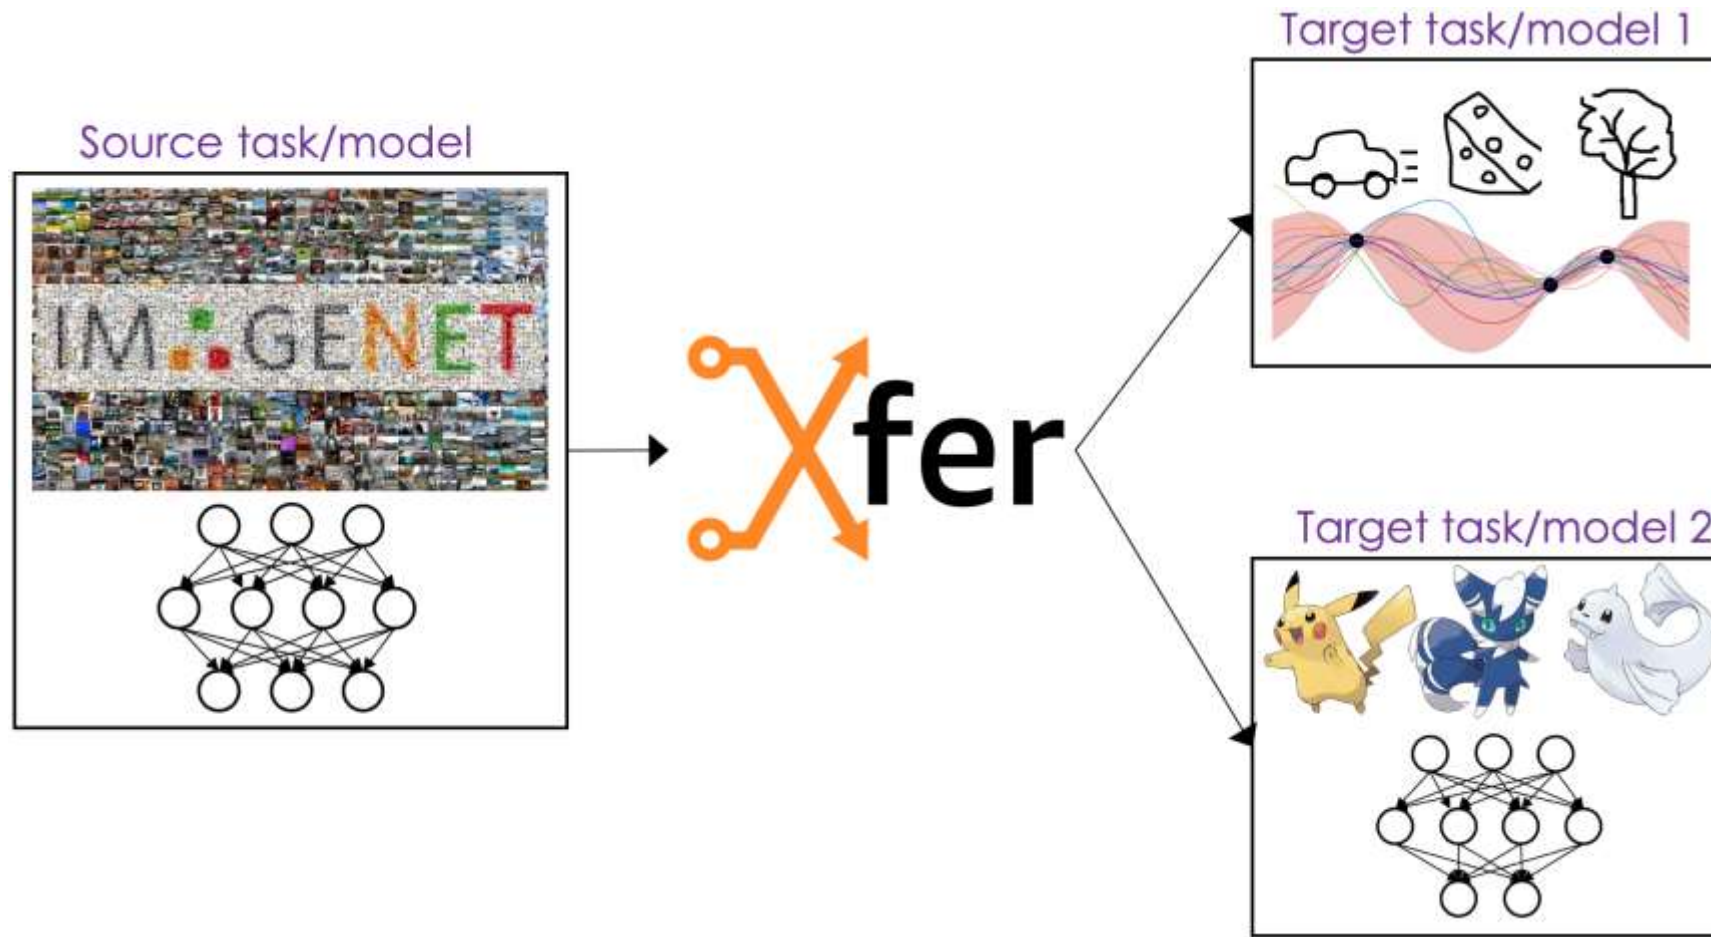


Transfer Learning


Use the **knowledge** gained while solving **one problem** and **applying** it to a **different** but related problem.

Start with a **pre-trained model** that are good at one task, lets you train far more **quickly** and with **less data** than if you were to train from scratch.


Transfer Learning





Faces






Cars

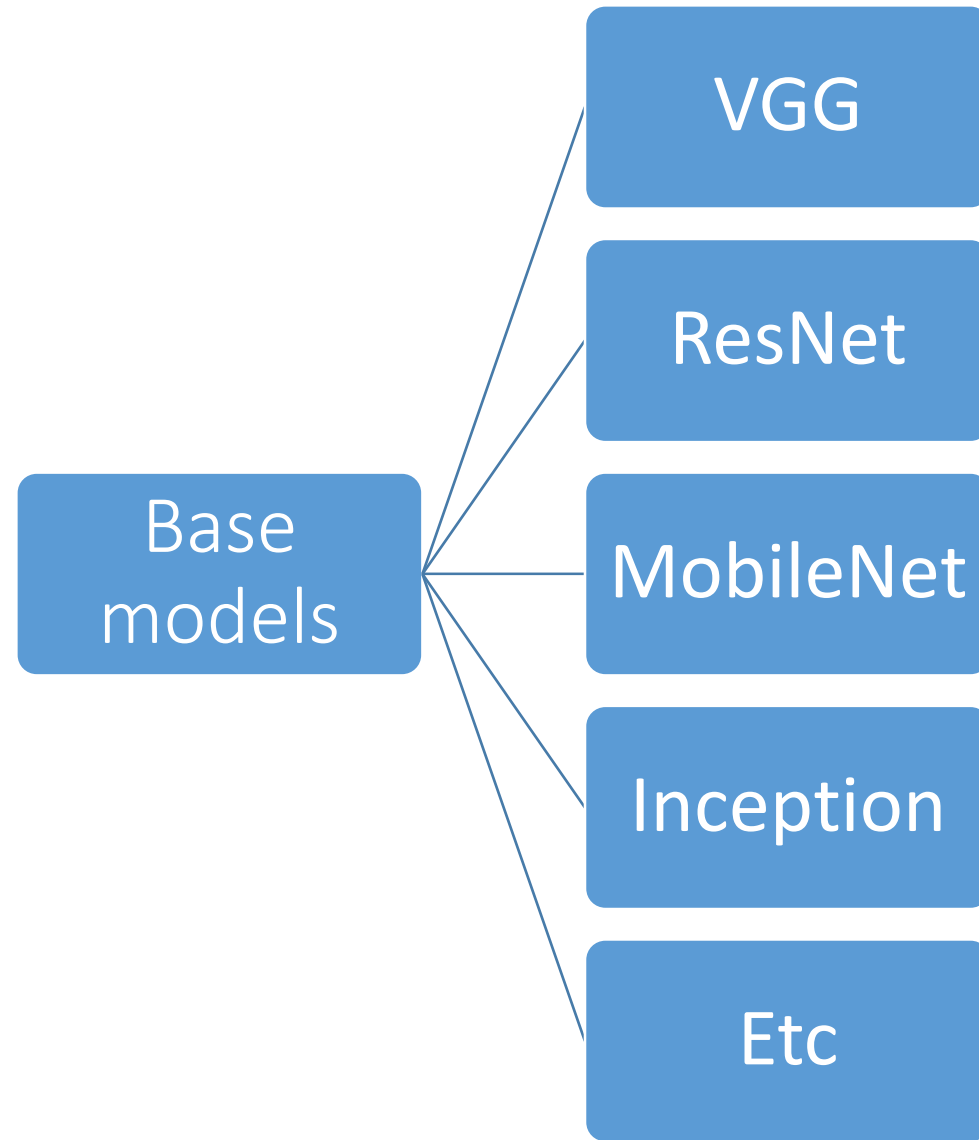


Elephants




Chairs





Transfer learning model



CHALLENGE

- Train the best classifier on the cats dogs dataset with transfer learning.
- Tune the base model and other hyperparameters for best results.
- With TL, you can get away training with very little data. But how little?

Lab 3: Teachable Machine


Customizable image classifier on webcam images

Teachable Machine

Train a computer to recognize your own images, sounds, & poses.

A fast, easy way to create machine learning models for your sites, apps, and more – no expertise or coding required.

[Get Started](#)



MobileNet Model

- **MobileNet: Efficient** Convolutional Neural Networks for Mobile Vision Applications([paper](#)).



Model Loaded

The MobileNet model labeled this as robin, American robin, *Turdus migratorius*, with a confidence of 0.99.



class name	probability	imagenet class id
Egyptian cat	0.478	285
tabby, tabby cat	0.300	281
tiger cat	0.167	282
remote control, remote	0.016	761
Siamese cat, Siamese	0.008	284

MobileNet Model

- **Small, fast, accurate.** A model that is trained on **ImageNet**.





Table 9. Smaller MobileNet Comparison to Popular Models

Model	ImageNet Accuracy	Million Mult-Adds	Million Parameters
0.50 MobileNet-160	60.2%	76	1.32
Squeezenet	57.5%	1700	1.25
AlexNet	57.2%	720	60


* Multiply-Accumulates (MACs) measures the number of fused Multiplication and Addition operations

ImageNet Dataset

ImageNet: a **dataset** of **millions** of images with labels for **1000** different classes of **objects**, like dogs, cats, and fruits.

ImageNet 2011 Fall Release (32326)

- plant, flora, plant life (4486)
- geological formation, formation (175)
- natural object (1112)
- sport, athletics (176)
- artifact, artefact (10504)
- fungus (308)
- person, individual, someone, somebody, mortal, soul (6978)
- animal, animate being, beast, brute, creature, fauna (3998)
- Misc (20400)



Transfer Learning

Use the knowledge gained while recognizing the **Imagenet classes** could apply when trying to recognize **user customized classes**.

Benefits:

- **Short** training time
- Needs very **little data**
- All in the **browser**.

Image from Webcam → [MobileNet](#)(Second to last layer) → [NN Classifier](#) → Output

(High-level semantic features of the image)

Introduction to Machine Learning & Deep Learning Concepts



ITXOTIC